



net: Go's DNS resolver fails with no such host #15434

New issue

Closed fraenkel opened this issue on Apr 26, 2016 · 9 comments



fraenkel commented on Apr 26, 2016

Contributor



Please answer these questions before submitting your issue. Thanks!

1. What version of Go are you using (`go version`)?
go 1.6.2
go version devel +093ac15 Mon Apr 25 06:00:15 2016 +0000 linux/amd64
2. What operating system and processor architecture are you using (`go env`)?
linux amd64
3. What did you do?
If possible, provide a recipe for reproducing the error.
A complete runnable program is good.
A link on play.golang.org is best.

<http://play.golang.org/p/O1SMk6JM1w>

1. What did you expect to see?
success
2. What did you see instead?
Get <http://registry-1.docker.io/v2>: dial tcp: lookup registry-1.docker.io on 10.155.248.190:53: no such host

When I set export GODEBUG=netdns=cgo+1, I do succeed but with the default or setting it to go, I always get the no such host.
I have 3 DNS entries in my resolv.conf. The second and third entries can resolve the name.
With strace, I see that the cgo trace shows a request to each DNS resolver, however with Go, it only tries the first and fails.

Assignees

No one assigned

Labels

None yet

Projects

None yet

Milestone

Go1.7

Notifications

Subscribe

You're not receiving notifications from this thread.

4 participants



mdempsky commented on Apr 26, 2016

Member



Can you share the output from running `dig registry-1.docker.io @10.155.248.190` on that machine?



fraenkel commented on Apr 26, 2016

Contributor



```
root@6515733c-cf49-4443-b7d9-adb073ea8885:~# dig registry-1.docker.io @10.155.248.190
```

```
; <<>> DiG 9.9.5-3ubuntu0.8-Ubuntu <<>> registry-1.docker.io @10.155.248.190
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 6235
;; flags: qr rd; QUERY: 1, ANSWER: 0, AUTHORITY: 0, ADDITIONAL: 1
;; WARNING: recursion requested but not available

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags::, udp: 2800
;; QUESTION SECTION:
;registry-1.docker.io. IN A

;; Query time: 1 msec
;; SERVER: 10.155.248.190#53(10.155.248.190)
;; WHEN: Mon Apr 25 17:05:55 UTC 2016
;; MSG SIZE rcvd: 49
```



mdempsky commented on Apr 26, 2016

Member



Interesting. With strace, can you tell if libc immediately tries the next nameserver, or does it continue waiting a bit longer before sending the next packet (e.g., calling poll(2) again)?

Looking at libresolv's source, I'm wondering if you're hitting this code:

```
if (anhp->rcode == NOERROR && anhp->ancount == 0
    && anhp->aa == 0 && anhp->ra == 0 && anhp->arcount == 0) {
    Dprintf(statp->options & RES_DEBUG,
        (stdout, "referred query:\n"),
        *thisansp,
        (*thisresplenp > *thisansszip)
        ? *thisansszip : *thisresplenp);
    goto next_ns;
}
```

except that according to the dig output, anhp->arcount (aka "ADDITIONAL") should be 1.



fraenkel commented on Apr 26, 2016

Contributor



Here is the relevant strace from both

go DNS

```
socket(PF_INET, SOCK_DGRAM|SOCK_CLOEXEC|SOCK_NONBLOCK, IPPROTO_IP) = 5
setsockopt(5, SOL_SOCKET, SO_BROADCAST, [1], 4) = 0
connect(5, {sa_family=AF_INET, sin_port=htons(53), sin_addr=inet_addr("10.155.248.190")}, 1)
epoll_ctl(4, EPOLL_CTL_ADD, 5, {EPOLLIN|EPOLLOUT|EPOLLRDHUP|EPOLLET, {u32=2021793344, u64=1}})
getsockname(5, {sa_family=AF_INET, sin_port=htons(50722), sin_addr=inet_addr("10.155.248.16")}, {sa_family=AF_INET, sin_port=htons(53), sin_addr=inet_addr("10.155.248.190")})
clock_gettime(CLOCK_REALTIME, {1461610393, 390514267}) = 0
clock_gettime(CLOCK_MONOTONIC, {5648, 84800638}) = 0
clock_gettime(CLOCK_MONOTONIC, {5648, 84874068}) = 0
clock_gettime(CLOCK_REALTIME, {1461610393, 390734946}) = 0
write(5, "\310-\1\0\0\1\0\0\0\0\0\0\0\nregistry-1\6docker\2i"... , 38) = 38
read(5, 0xc8200ee000, 512) = -1 EAGAIN (Resource temporarily unavailable)
futex(0x704bd0, FUTEX_WAIT, 0, NULLGet http://registry-1.docker.io/v2: dial tcp: lookup reg
```

cgo DNS

```
socket(PF_INET, SOCK_DGRAM|SOCK_NONBLOCK, IPPROTO_IP) = 3
connect(3, {sa_family=AF_INET, sin_port=htons(53), sin_addr=inet_addr("10.155.248.190")}, 1)
gettimeofday({1461610383, 13246}, NULL) = 0
poll([{fd=3, events=POLLOUT}], 1, 0) = 1 ([{fd=3, revents=POLLOUT}])
sendmmsg(3, [{msg_name(0)=NULL, msg_iov(1)=[{"\324^\1\0\0\1\0\0\0\0\0\0\nregistry-1\6docke"}], 1, 0) = 1 ([{fd=3, revents=POLLIN}])
poll([{fd=3, events=POLLIN}], 1, 5000) = 1 ([{fd=3, revents=POLLIN}])
ioctl(3, FIONREAD, [38]) = 0
recvfrom(3, "\364\244\201\0\0\1\0\0\0\0\0\0\nregistry-1\6docker\2i"... , 2048, 0, {sa_family=AF_INET, sin_port=htons(53), sin_addr=inet_addr("10.0.80.11")}, 16) = 2048
gettimeofday({1461610383, 14478}, NULL) = 0
poll([{fd=3, events=POLLOUT}], 1, 0) = 1 ([{fd=3, revents=POLLOUT}])
sendmmsg(3, [{msg_name(0)=NULL, msg_iov(1)=[{"\324^\1\0\0\1\0\0\0\0\0\0\nregistry-1\6docke"}], 1, 0) = 1 ([{fd=3, revents=POLLIN}])
poll([{fd=3, events=POLLIN}], 1, 3000) = 1 ([{fd=3, revents=POLLIN}])
ioctl(3, FIONREAD, [125]) = 0
recvfrom(3, "\324^\201\200\0\1\0\0\0\0\0\0\nregistry-1\6docker\2i"... , 2048, 0, {sa_family=AF_INET, sin_port=htons(53), sin_addr=inet_addr("10.0.80.11")}, 16) = 2048
gettimeofday({1461610383, 15315}, NULL) = 0
poll([{fd=3, events=POLLIN}], 1, 2999) = 1 ([{fd=3, revents=POLLIN}])
ioctl(3, FIONREAD, [152]) = 0
recvfrom(3, "\364\244\201\200\0\1\0\0\0\0\0\nregistry-1\6docker\2i"... , 65536, 0, {sa_family=AF_INET, sin_port=htons(53), sin_addr=inet_addr("10.0.80.11")}, 16) = 65536
close(3) = 0
```



mdempsky commented on Apr 26, 2016

Member



Thanks! Looking at the DNS response header from 10.155.248.190 in the cgo strace output, it's clear that the response *does* actually have arcount==0, which explains why libresolv proceeds to try the next name server.

Replicating libresolv's behavior here seems easy, but feels kinda gross as it doesn't seem to be prescribed by the RFCs at all.

FYI though, your netgo strace output looks incomplete though, like you're only tracing one of the OS threads. In particular, it's missing the `write` call that actually produced the "no such host" output, and it's probably missing the `recvmsg` call that actually read the response DNS packet; the one in the output above failed with `EAGAIN`.

I usually use `strace -f` to trace all children.



fraenkel commented on Apr 26, 2016

Contributor +

I have a more more complete go trace with some editorials to trim the unwanted stuff.

I am noticing that we connect to 10.155.248.190 twice. Not sure if that is expected.

```
[pid 7617] socket(PF_INET, SOCK_DGRAM|SOCK_CLOEXEC|SOCK_NONBLOCK, IPPROTO_IP <unfinished .
[pid 7621] futex(0x703e20, FUTEX_WAIT, 2, NULL <unfinished ...>
[pid 7617] <... socket resumed> ) = 3
[pid 7617] setsockopt(3, SOL_SOCKET, SO_BROADCAST, [1], 4 <unfinished ...>
[pid 7617] <... setsockopt resumed> ) = 0
[pid 7619] futex(0x703e20, FUTEX_WAKE, 1 <unfinished ...>
[pid 7617] connect(3, {sa_family=AF_INET, sin_port=htons(53), sin_addr=inet_addr("10.155.2
[pid 7621] <... futex resumed> ) = 0
[pid 7617] <... connect resumed> ) = 0
[pid 7619] <... futex resumed> ) = 1
[pid 7617] epoll_create1(EPOLL_CLOEXEC <unfinished ...>
[pid 7621] futex(0x703e20, FUTEX_WAKE, 1 <unfinished ...>
[pid 7617] <... epoll_create1 resumed> ) = 4
[pid 7621] <... futex resumed> ) = 0
[pid 7617] epoll_ctl(4, EPOLL_CTL_ADD, 3, {EPOLLIN|EPOLLOUT|EPOLLRDHUP|EPOLLET, {u32=13877
[pid 7617] <... epoll_ctl resumed> ) = 0
[pid 7621] clock_gettime(CLOCK_REALTIME, <unfinished ...>
[pid 7617] getsockname(3, <unfinished ...>
[pid 7617] <... getsockname resumed> {sa_family=AF_INET, sin_port=htons(59996), sin_addr=i
[pid 7617] getpeername(3, <unfinished ...>
[pid 7617] <... getpeername resumed> {sa_family=AF_INET, sin_port=htons(53), sin_addr=inet
[pid 7621] socket(PF_INET, SOCK_DGRAM|SOCK_CLOEXEC|SOCK_NONBLOCK, IPPROTO_IP <unfinished .
[pid 7621] <... socket resumed> ) = 5
[pid 7621] setsockopt(5, SOL_SOCKET, SO_BROADCAST, [1], 4 <unfinished ...>
[pid 7621] <... setsockopt resumed> ) = 0
[pid 7621] connect(5, {sa_family=AF_INET, sin_port=htons(53), sin_addr=inet_addr("10.155.2
[pid 7621] <... connect resumed> ) = 0
[pid 7617] futex(0x703e38, FUTEX_WAKE, 1 <unfinished ...>
[pid 7621] epoll_ctl(4, EPOLL_CTL_ADD, 5, {EPOLLIN|EPOLLOUT|EPOLLRDHUP|EPOLLET, {u32=13877
[pid 7617] <... futex resumed> ) = 0
[pid 7621] <... epoll_ctl resumed> ) = 0
[pid 7621] getsockname(5, <unfinished ...>
[pid 7621] <... getsockname resumed> {sa_family=AF_INET, sin_port=htons(43899), sin_addr=i
[pid 7617] write(3, "f\273\1\0\0\1\0\0\0\0\0\nregistry-1\6docker\2i"... , 38 <unfinished
[pid 7621] getpeername(5, <unfinished ...>
[pid 7619] futex(0x703e38, FUTEX_WAIT, 0, {4, 999648197} <unfinished ...>
[pid 7617] <... write resumed> ) = 38
[pid 7621] <... getpeername resumed> {sa_family=AF_INET, sin_port=htons(53), sin_addr=inet
[pid 7619] <... futex resumed> ) = -1 EAGAIN (Resource temporarily unavailable)
[pid 7617] read(3, <unfinished ...>
[pid 7617] <... read resumed> 0xc8200e8000, 512) = -1 EAGAIN (Resource temporarily unavail
[pid 7617] epoll_wait(4, <unfinished ...>
[pid 7617] <... epoll_wait resumed> {{EPOLLOUT, {u32=1387773736, u64=139927127310120}}, {E
[pid 7617] epoll_wait(4, <unfinished ...>
[pid 7621] sched_yield( <unfinished ...>
[pid 7621] <... sched_yield resumed> ) = 0
[pid 7621] futex(0x703e20, FUTEX_WAKE, 1 <unfinished ...>
[pid 7619] futex(0x703e38, FUTEX_WAIT, 0, {4, 998688799} <unfinished ...>
[pid 7621] <... futex resumed> ) = 0
[pid 7618] <... select resumed> ) = 0 (Timeout)
[pid 7621] write(5, "4\274\1\0\0\1\0\0\0\0\0\nregistry-1\6docker\2i"... , 38 <unfinished
[pid 7621] <... write resumed> ) = 38
[pid 7617] <... epoll_wait resumed> {{EPOLLOUT, {u32=1387773544, u64=139927127309928}}, 1
[pid 7621] read(5, <unfinished ...>
[pid 7617] epoll_wait(4, <unfinished ...>
[pid 7618] select(0, NULL, NULL, NULL, {0, 20} <unfinished ...>
[pid 7617] <... epoll_wait resumed> {{EPOLLIN|EPOLLOUT, {u32=1387773736, u64=1399271273101
[pid 7621] <... read resumed> 0xc8200ee000, 512) = -1 EAGAIN (Resource temporarily unavail
[pid 7621] futex(0xc82005a110, FUTEX_WAIT, 0, NULL <unfinished ...>
[pid 7617] read(3, "f\273\201\0\0\1\0\0\0\0\0\nregistry-1\6docker\2i"... , 512) = 38
[pid 7617] epoll_ctl(4, EPOLL_CTL_DEL, 3, {0, {u32=0, u64=0}} <unfinished ...>
[pid 7618] <... select resumed> ) = 0 (Timeout)
[pid 7617] <... epoll_ctl resumed> ) = 0
[pid 7617] close(3 <unfinished ...>
[pid 7617] <... close resumed> ) = 0
[pid 7617] futex(0xc82005a110, FUTEX_WAKE, 1 <unfinished ...>
[pid 7617] <... futex resumed> ) = 1
[pid 7621] <... futex resumed> ) = 0
[pid 7617] epoll_wait(4, <unfinished ...>
[pid 7621] epoll_wait(4, <unfinished ...>
```

```
[pid 7617] <... epoll_wait resumed> {{EPOLLIN|EPOLLOUT, {u32=1387773544, u64=1399271273099
[pid 7621] <... epoll_wait resumed> {}, 128, 0) = 0
[pid 7617] read(5, "4\274\201\0\0\1\0\0\0\0\0\nregistry-1\6docker\2i"... , 512) = 38
[pid 7621] epoll_wait(4, <unfinished ...>
[pid 7618] select(0, NULL, NULL, NULL, {0, 20} <unfinished ...>
[pid 7617] epoll_ctl(4, EPOLL_CTL_DEL, 5, {0, {u32=0, u64=0}}) = 0
[pid 7617] close(5) = 0
[pid 7617] futex(0xc82002ed10, FUTEX_WAKE, 1) = 1
[pid 7620] <... futex resumed> ) = 0
[pid 7618] <... select resumed> ) = 0 (Timeout)
[pid 7617] write(2, "Get http://registry-1.docker.io/"..., 108 <unfinished ...>
Get http://registry-1.docker.io/v2: dial tcp: lookup registry-1.docker.io on 10.155.248.190
```



mdempsky commented on Apr 26, 2016

Member



@fraenkel Thanks, that looks more like what I would expect.

I suspect the two connects are because we're issuing both A and AAAA queries, but the strace output is truncating the DNS packets so I can't see the QTYPE field.

Anyway, I just uploaded <https://go-review.googlesource.com/#/c/22428/>, which I believe should fix the issue by replicating libresolv's behavior. Do you mind testing if it fixes your problem?



fraenkel commented on Apr 26, 2016

Contributor



It works. Thanks for the quick turnaround.



gopherbot commented on Apr 26, 2016



CL <https://golang.org/cl/22428> mentions this issue.

gopherbot closed this in [98b99d5](#) on Apr 27, 2016

★ mdempsky referenced this issue on Jun 30, 2016

net: LookupHost only using the first configure DNS server in /etc/resolv.conf on Linux #16215

Closed

mikioh changed the title from **Go's DNS resolver fails with no such host** to **net: Go's DNS resolver fails with no such host** on Oct 19, 2016

mikioh added this to the **Go1.7** milestone on Oct 19, 2016



Write Preview

AA B i



Leave a comment

Attach files by dragging & dropping, [selecting them](#), or pasting from the clipboard.

Styling with Markdown is supported

Comment

