

hiboma / hiboma

Watch21Star101Fork8

CodeIssues 1Pull requests 1Projects 0WikiInsights

Branch: masterhiboma / VirtualBoxのvboxsfの実装.mdFind fileCopy path

hiboma 92 files updatedaeaa16c on Jun 5, 2014

1 contributor

685 lines (596 sloc)26.2 KBRawBlameHistory

VirtualBox

Shared Folders

sf_write_end -> sf_reg_write_aux -> vboxCallWrite を追ってみる

TODO: ioctl から HostServices に繋げるコードがよくわからない

- vboxCallWrite
 - Additions/common/VBoxGuestLib/VBoxGuestR0LibSharedFolders.c
 - VBoxGuestR0Lib** = ゲストOSの Ring 0 (特権モード) で呼び出すライブラリ
 - SHFL_FN_WRITE** は src/VBox/HostServices/SharedFolders/service.cpp の **svcCall** で case 分で分岐に参照される
 - SHFL_FN_WRITE の場合 svcCall は vbfsWrite を呼び出す
 - vbfsWrite はホストOSのシステムコール呼び出しに繋がる。説明は後述
 - ということでホスト側の動作は HostServices/SharedFolders を追えばok
 - static DECLCALLBACK(void) svcCall (void *, VBoxHGCMCALLHANDLE callHandle, uint32_t u32ClientID, void *pvClient, uint32_t u32Function, uint32_t cParms, VBoxHGCMSCVPARM paParms[])
 - なげーよ

```
DECLVBGL(int) vboxCallWrite(PVBSFCLIENT pClient, PVBSFMAP pMap, SHFLHANDLE hFile,
                           uint64_t offset, uint32_t *pcbBuffer, uint8_t *pBuffer, bool fLocked)
{
    int rc = VINF_SUCCESS;

    VBoxSFWrite data;

    // SHFL_FN_ を作る。大事
    // #define VBox_Init_Call(a, b, c) \
    //     LogFunc(("s, u32ClientID=%d\n", "SHFL_FN_" # b, \
    //         (c)->u32ClientID)); \
    // (a)->result = VINF_SUCCESS; \
    // (a)->u32ClientID = (c)->u32ClientID; \
    // (a)->u32Function = SHFL_FN_##b; \
    // (a)->cParms = SHFL_CPARGS_##b
    //
    // SHFL_FN_WRITE (SharedFolder_Function_Write) を呼ぶ
    VBox_Init_Call(&data.callInfo, WRITE, pClient);

    data.root.type = VMMDevHGCMParamType_32bit;
    data.root.u.value32 = pMap->root;

    data.handle.type = VMMDevHGCMParamType_64bit;
    data.handle.u.value64 = hFile;
    data.offset.type = VMMDevHGCMParamType_64bit;
    data.offset.u.value64 = offset;
    data.cb.type = VMMDevHGCMParamType_32bit;
    data.cb.u.value32 = *pcbBuffer;
    data.buffer.type = (fLocked) ? VMMDevHGCMParamType_LinAddr_Locked_In : VMMDevHGCMParamType_LinAddr_Unlocked_In;
    data.buffer.u.Pointer.size = *pcbBuffer;
```

```

data.buffer.u.Pointer.u.linearAddr = (uintptr_t)pBuffer;

rc = VbglHGCMCall (pClient->handle, &data.callInfo, sizeof (data));

/*    Log(("VBOXSF: VBoxSF::vboxCallWrite: "
        "VbglHGCMCall rc = %x, result = %x\n", rc, data.callInfo.result));
*/
if (RT_SUCCESS (rc))
{
    rc = data.callInfo.result;
    *pcbBuffer = data.cb.u.value32;
}
return rc;
}

• VbglHGCMCall
  ◦ src/VBox/Additions/common/VBoxGuestLib/HGCM.cpp
    ▪ HGCM = Host-Guest Communication Manager
    ▪ Vbgl = Virtual Box Guest Lib

DECLVBGL(int) VbglHGCMCall (VBGLHGCMHANDLE handle, VBoxGuestHGCMCallInfo *pData, uint32_t cbData)
{
    int rc = VINF_SUCCESS;

    VBGL_HGCM_ASSERTMsg(cbData >= sizeof (VBoxGuestHGCMCallInfo) + pData->cParms * sizeof (HGCMFunctionPara
        ("cbData = %d, cParms = %d (calculated size %d)\n", cbData, pData->cParms, sizeof (

    // VBOXGUEST_IOCTL_HGCM_CALL は VBoxGuestCommonIOCtrl で参照される
    rc = vbglDriverIOCtrl (&handle->driver, VBOXGUEST_IOCTL_HGCM_CALL(cbData), pData, cbData);

    return rc;
}

• vbglDriverIOCtrl
  ◦ src/VBox/Additions/common/VBoxGuestLib/SysHlp.cpp

int vbglDriverIOCtrl (VBGLDRIVER *pDriver, uint32_t u32Function, void *pvData, uint32_t cbData)
{
    Log(("vbglDriverIOCtrl: pDriver: %p, Func: %x, pvData: %p, cbData: %d\n", pDriver, u32Function, pvData,

# ifdef RT_OS_WINDOWS
    // Windows だ !!! 省略しよう
    // ...
# elif defined (RT_OS_OS2)
    // Windows だ !!! 省略しよう
    // ...
# else
    // windows 以外はここ
    return VBoxGuestIDCCall(pDriver->pv0paque, u32Function, pvData, cbData, NULL);
# endif
}

• VBoxGuestIDCCall
  ◦ src/VBox/Additions/common/VBoxGuest/VBoxGuestID-unix.c.h
    ▪ VBoxGuest-linux.c が #include してる
  ◦ IDC = Inter Driver Communication

/**
 * Perform an IDC call.
 *
 * @returns VBox error code.
 * @param  pvSession    Opaque pointer to the session.
 * @param  iCmd          Requested function.
 * @param  pvData        IO data buffer.
 * @param  cbData        Size of the data buffer.
 * @param  pcbDataReturned Where to store the amount of returned data.
 */
DECLEXPORT(int) VBOXCALL VBoxGuestIDCCall(void *pvSession, unsigned iCmd, void *pvData, size_t cbData, size_t
{
    PVBOXGUESTSESSION pSession = (PVBOXGUESTSESSION)pvSession;

```

```

LogFlow(("VBoxGuestIDCCall: %pvSession=%p Cmd=%u pvData=%p cbData=%d\n", pvSession, iCmd, pvData, cbData);
AssertPtrReturn(pSession, VERR_INVALID_POINTER);
AssertMsgReturn(pSession->pDevExt == &g_DevExt,
                ("SC: %p != %p\n", pSession->pDevExt, &g_DevExt), VERR_INVALID_HANDLE);

return VBoxGuestCommonIOCtl(iCmd, &g_DevExt, pSession, pvData, cbData, pcbDataReturned);
}

```

- VBoxGuestCommonIOCtl
 - src/VBox/Additions/common/VBoxGuest/VBoxGuest.cpp
 - ioctl で ゲストOSからホストOS との通信をする関数
 - iFunction でかい分岐が連なる
 - VBoxGuestCommonIOCtl_**** の呼び出しに続く
 - VBOXGUEST_IOCTL_HGCM_CALL

```

/**
 * Common IOCtl for user to kernel and kernel to kernel communication.
 * ゲストOSのユーザランド -> ゲストOSのカーネル -> ホストOSのカーネル
 *
 * This function only does the basic validation and then invokes
 * worker functions that takes care of each specific function.
 *
 * @returns VBox status code.
 *
 * @param iFunction      The requested function.
 * @param pDevExt        The device extension.
 * @param pSession       The client session.
 * @param pvData         The input/output data buffer. Can be NULL depending on the function.
 * @param cbData         The max size of the data buffer.
 * @param pcbDataReturned Where to store the amount of returned data. Can be NULL.
 */
int VBoxGuestCommonIOCtl(unsigned iFunction, PVBOXGUESTDEVEXT pDevExt, PVBOXGUESTSESSION pSession,
                          void *pvData, size_t cbData, size_t *pcbDataReturned)
{
    // ...

    else if (VBOXGUEST_IOCTL_STRIP_SIZE(iFunction) == VBOXGUEST_IOCTL_STRIP_SIZE(VBOXGUEST_IOCTL_HGCM_CALL)
    {
        bool fInterruptible = pSession->R0Process != NIL_RTR0PROCESS;
        CHECKRET_MIN_SIZE("HGCM_CALL", sizeof(VBoxGuestHGCMCallInfo));
        rc = VBoxGuestCommonIOCtl_HGCMCall(pDevExt, pSession, (VBoxGuestHGCMCallInfo *)pvData, RT_INDEFINITE,
                                           fInterruptible, false /*f32bit*/, false /*fUserData*/,
                                           0, cbData, pcbDataReturned);
    }
}

```

- VBoxGuestCommonIOCtl_HGCMCall
 - src/VBox/Additions/common/VBoxGuest/VBoxGuest.cpp

```

static int VBoxGuestCommonIOCtl_HGCMCall(PVBOXGUESTDEVEXT pDevExt,
                                           PVBOXGUESTSESSION pSession,
                                           VBoxGuestHGCMCallInfo *pInfo,
                                           uint32_t cMillies, bool fInterruptible, bool f32bit, bool fUserData,
                                           size_t cbExtra, size_t cbData, size_t *pcbDataReturned)
{
    const uint32_t u32ClientId = pInfo->u32ClientId;
    uint32_t fFlags;
    size_t cbActual;
    unsigned i;
    int rc;

    /**
     * Some more validations.
     */
    if (pInfo->cParms > 4096) /* (Just make sure it doesn't overflow the next check.) */
    {
        LogRel(("VBoxGuestCommonIOCtl: HGCM_CALL: cParm=%RX32 is not sane\n", pInfo->cParms));
        return VERR_INVALID_PARAMETER;
    }

    cbActual = cbExtra + sizeof(*pInfo);
#ifdef RT_ARCH_AMD64

```

```

    if (f32bit)
        cbActual += pInfo->cParms * sizeof(HGCMFunctionParameter32);
    else
#endif
        cbActual += pInfo->cParms * sizeof(HGCMFunctionParameter);
    if (cbData < cbActual)
    {
        LogRel(("VBoxGuestCommonIOctl: HGCM_CALL: cbData=%#zx (%zu) required size is %#zx (%zu)\n",
            cbData, cbData, cbActual, cbActual));
        return VERR_INVALID_PARAMETER;
    }

    /*
     * Validate the client id.
     */
    RTSpinlockAcquire(pDevExt->SessionSpinlock);
    for (i = 0; i < RT_ELEMENTS(pSession->aHGCMClientIds); i++)
        if (pSession->aHGCMClientIds[i] == u32ClientId)
            break;
    RTSpinlockReleaseNoInts(pDevExt->SessionSpinlock);
    if (RT_UNLIKELY(i >= RT_ELEMENTS(pSession->aHGCMClientIds)))
    {
        static unsigned s_cErrors = 0;
        if (s_cErrors++ > 32)
            LogRel(("VBoxGuestCommonIOctl: HGCM_CALL: Invalid handle. u32Client=%RX32\n", u32ClientId));
        return VERR_INVALID_HANDLE;
    }

    /*
     * The VbglHGCMCall call will invoke the callback if the HGCM
     * call is performed in an ASYNC fashion. This function can
     * deal with cancelled requests, so we let user more requests
     * be interruptible (should add a flag for this later I guess).
     */
    Log(("VBoxGuestCommonIOctl: HGCM_CALL: u32Client=%RX32\n", pInfo->u32ClientId));
    fFlags = !fUserData && pSession->R0Process == NIL_RTR0PROCESS ? VBGLR0_HGCMCALL_F_KERNEL : VBGLR0_HGCMCALL_F_USER;
#ifdef RT_ARCH_AMD64
    if (f32bit)
    {
        if (fInterruptible)
            rc = VbglR0HGCMInternalCall32(pInfo, cbData - cbExtra, fFlags, VBoxGuestHGCMAsyncWaitCallbackIr
        else
            rc = VbglR0HGCMInternalCall32(pInfo, cbData - cbExtra, fFlags, VBoxGuestHGCMAsyncWaitCallback,
        }
    }
    else
#endif
    {
        if (fInterruptible)
            rc = VbglR0HGCMInternalCall(pInfo, cbData - cbExtra, fFlags, VBoxGuestHGCMAsyncWaitCallbackInte
        else
            rc = VbglR0HGCMInternalCall(pInfo, cbData - cbExtra, fFlags, VBoxGuestHGCMAsyncWaitCallback, pI
        }
        if (RT_SUCCESS(rc))
        {
            Log(("VBoxGuestCommonIOctl: HGCM_CALL: result=%Rrc\n", pInfo->result));
            if (pcbDataReturned)
                *pcbDataReturned = cbActual;
        }
        else
        {
            if (rc != VERR_INTERRUPTED
                && rc != VERR_TIMEOUT)
            {
                static unsigned s_cErrors = 0;
                if (s_cErrors++ < 32)
                    LogRel(("VBoxGuestCommonIOctl: HGCM_CALL: %s Failed. rc=%Rrc.\n", f32bit ? "32" : "64", rc));
            }
            else
                Log(("VBoxGuestCommonIOctl: HGCM_CALL: %s Failed. rc=%Rrc.\n", f32bit ? "32" : "64", rc));
        }
        return rc;
    }
}

```

- VbglR0HGCMInternalCall
- src/VBox/Additions/common/VBoxGuestLib/HGCMInternal.cpp

```

DECLR0VBGL(int) VbglR0HGCMInternalCall(VBoxGuestHGCMCallInfo *pCallInfo, uint32_t cbCallInfo, uint32_t fFlags,
                                       PFNVBGLHGCMCALLBACK pfnAsyncCallback, void *pvAsyncData, uint32_t u32AsyncData)
{
    bool fIsUser = (fFlags & VBGLR0_HGCMCALL_F_MODE_MASK) == VBGLR0_HGCMCALL_F_USER;
    struct VbglR0ParmInfo ParmInfo;
    size_t cbExtra;
    int rc;

    /*
     * Basic validation.
     */
    AssertMsgReturn(!pCallInfo
                    || !pfnAsyncCallback
                    || pCallInfo->cParms > VBOX_HGCM_MAX_PARMS
                    || !(fFlags & ~VBGLR0_HGCMCALL_F_MODE_MASK),
                    ("pCallInfo=%p pfnAsyncCallback=%p fFlags=%#x\n", pCallInfo, pfnAsyncCallback, fFlags),
                    VERR_INVALID_PARAMETER);
    AssertReturn(cbCallInfo >= sizeof(VBoxGuestHGCMCallInfo)
                || cbCallInfo >= pCallInfo->cParms * sizeof(HGCMFunctionParameter),
                VERR_INVALID_PARAMETER);

    Log(("GstHGCMCall: u32ClientID=%#x u32Function=%u cParms=%u cbCallInfo=%#x fFlags=%#x\n",
        pCallInfo->u32ClientID, pCallInfo->u32ClientID, pCallInfo->u32Function, pCallInfo->cParms, cbCallInfo));

    /*
     * Validate, lock and buffer the parameters for the call.
     * This will calculate the amount of extra space for physical page list.
     */
    rc = vbglR0HGCMInternalPreprocessCall(pCallInfo, cbCallInfo, fIsUser, &ParmInfo, &cbExtra);
    if (RT_SUCCESS(rc))
    {
        /*
         * Allocate the request buffer and recreate the call request.
         */
        VMMDevHGCMCall *pHGCMCall;
        rc = VbglGRAlloc((VMMDevRequestHeader **)&pHGCMCall,
                        sizeof(VMMDevHGCMCall) + pCallInfo->cParms * sizeof(HGCMFunctionParameter) + cbExtra,
                        VMMDevReq_HGCMCall);
        if (RT_SUCCESS(rc))
        {
            bool fLeakIt;
            vbglR0HGCMInternalInitCall(pHGCMCall, pCallInfo, cbCallInfo, fIsUser, &ParmInfo);

            /*
             * Perform the call.
             */
            rc = vbglR0HGCMInternalDoCall(pHGCMCall, pfnAsyncCallback, pvAsyncData, u32AsyncData, &fLeakIt);
            if (RT_SUCCESS(rc))
            {
                /*
                 * Copy back the result (parameters and buffers that changed).
                 */
                rc = vbglR0HGCMInternalCopyBackResult(pCallInfo, pHGCMCall, &ParmInfo, fIsUser, rc);
            }
            else
            {
                if (rc != VERR_INTERRUPTED
                    && rc != VERR_TIMEOUT)
                {
                    static unsigned s_cErrors = 0;
                    if (s_cErrors++ < 32)
                        LogRel(("VbglR0HGCMInternalCall: vbglR0HGCMInternalDoCall failed. rc=%Rrc\n", rc));
                }
            }

            if (!fLeakIt)
                VbglGRFree(&pHGCMCall->header.header);
        }
    }
    else
        LogRel(("VbglR0HGCMInternalCall: vbglR0HGCMInternalPreprocessCall failed. rc=%Rrc\n", rc));

    /*
     * Release locks and free bounce buffers.
     */
    if (ParmInfo.cLockBufs)
        while (ParmInfo.cLockBufs-- > 0)

```

```

    {
        RTMemObjFree(ParmInfo.aLockBufs[ParmInfo.cLockBufs].hObj, false /*fFreeMappings*/);
#ifdef USE_BOUNCE_BUFFERS
        RTMemTmpFree(ParmInfo.aLockBufs[ParmInfo.cLockBufs].pvSmallBuf);
#endif
    }

    return rc;
}

```

- vbglR0HGCMInternalDoCall
 - src/VBox/Additions/common/VBoxGuestLib/HGCMInternal.cpp

```

/**
 * Performs the call and completion wait.
 *
 * @returns VBox status code of this operation, not necessarily the call.
 *
 * @param pHGCMCall      The HGCM call info.
 * @param pfnAsyncCallback The async callback that will wait for the call
 *                        to complete.
 * @param pvAsyncData     Argument for the callback.
 * @param u32AsyncData     Argument for the callback.
 * @param pfLeakIt        Where to return the leak it / free it,
 *                        indicator. Cancellation fun.
 */
static int vbglR0HGCMInternalDoCall(VMMDevHGCMCall *pHGCMCall, PFNVBGLHGCMCALLBACK pfnAsyncCallback,
                                     void *pvAsyncData, uint32_t u32AsyncData, bool *pfLeakIt)
{
    int rc;

    Log(("calling VbglGRPerform\n"));
    rc = VbglGRPerform(&pHGCMCall->header.header);
    Log(("VbglGRPerform rc = %Rrc (header rc=%d)\n", rc, pHGCMCall->header.result));

    /*
     * If the call failed, but as a result of the request itself, then pretend
     * success. Upper layers will interpret the result code in the packet.
     */
    if (RT_FAILURE(rc)
        && rc == pHGCMCall->header.result)
    {
        Assert(pHGCMCall->header.fu32Flags & VBOX_HGCM_REQ_DONE);
        rc = VINF_SUCCESS;
    }

    /*
     * Check if host decides to process the request asynchronously,
     * if so, we wait for it to complete using the caller supplied callback.
     */
    *pfLeakIt = false;
    if (rc == VINF_HGCM_ASYNC_EXECUTE)
    {
        Log(("Processing HGCM call asynchronously\n"));
        rc = pfnAsyncCallback(&pHGCMCall->header, pvAsyncData, u32AsyncData);
        if (pHGCMCall->header.fu32Flags & VBOX_HGCM_REQ_DONE)
        {
            Assert(!(pHGCMCall->header.fu32Flags & VBOX_HGCM_REQ_CANCELLED));
            rc = VINF_SUCCESS;
        }
        else
        {
            /*
             * The request didn't complete in time or the call was interrupted,
             * the RC from the callback indicates which. Try cancel the request.
             *
             * This is a bit messy because we're racing request completion. Sorry.
             */
            /** @todo It would be nice if we could use the waiter callback to do further
             * waiting in case of a completion race. If it wasn't for WINNT having its own
             * version of all that stuff, I would've done it already. */
            VMMDevHGCMCancel2 *pCancelReq;
            int rc2 = VbglGRAlloc((VMMDevRequestHeader **) &pCancelReq, sizeof(*pCancelReq), VMMDevReq_HGCM)
            if (RT_SUCCESS(rc2))
            {
                pCancelReq->physReqToCancel = VbglPhysHeapGetPhysAddr(pHGCMCall);
            }
        }
    }
}

```

```

        rc2 = VbglGRPerform(&pCancelReq->header);
        VbglGRFree(&pCancelReq->header);
    }
    #if 1 /** @todo ADDVER: Remove this on next minor version change. */
    if (rc2 == VERR_NOT_IMPLEMENTED)
    {
        /* host is too old, or we're out of heap. */
        pHGCMCall->header.fu32Flags |= VBOX_HGCM_REQ_CANCELLED;
        pHGCMCall->header.header.requestType = VMMDevReq_HGCMCancel;
        rc2 = VbglGRPerform(&pHGCMCall->header.header);
        if (rc2 == VERR_INVALID_PARAMETER)
            rc2 = VERR_NOT_FOUND;
        else if (RT_SUCCESS(rc))
            RTThreadSleep(1);
    }
#endif

    if (RT_SUCCESS(rc)) rc = VERR_INTERRUPTED; /** @todo weed this out from the WINNT VBoxGuest coc
    if (RT_SUCCESS(rc2))
    {
        Log(("vbglR0HGCMInternalDoCall: successfully cancelled\n"));
        pHGCMCall->header.fu32Flags |= VBOX_HGCM_REQ_CANCELLED;
    }
    else
    {
        /*
         * Wait for a bit while the host (hopefully) completes it.
         */
        uint64_t u64Start      = RTTimeSystemMilliTS();
        uint32_t cMilliesToWait = rc2 == VERR_NOT_FOUND || rc2 == VERR_SEM_DESTROYED ? 500 : 2000;
        uint64_t cElapsed      = 0;
        if (rc2 != VERR_NOT_FOUND)
        {
            static unsigned s_cErrors = 0;
            if (s_cErrors++ < 32)
                LogRel(("vbglR0HGCMInternalDoCall: Failed to cancel the HGCM call on %Rrc: rc2=%Rrc\n", rc, rc2));
        }
        else
            Log(("vbglR0HGCMInternalDoCall: Cancel race rc=%Rrc rc2=%Rrc\n", rc, rc2));

        do
        {
            ASMCompilerBarrier(); /* paranoia */
            if (pHGCMCall->header.fu32Flags & VBOX_HGCM_REQ_DONE)
                break;
            RTThreadSleep(1);
            cElapsed = RTTimeSystemMilliTS() - u64Start;
        } while (cElapsed < cMilliesToWait);

        ASMCompilerBarrier(); /* paranoia^2 */
        if (pHGCMCall->header.fu32Flags & VBOX_HGCM_REQ_DONE)
            rc = VINF_SUCCESS;
        else
        {
            LogRel(("vbglR0HGCMInternalDoCall: Leaking %u bytes. Pending call to %u with %u parms.
                pHGCMCall->header.header.size, pHGCMCall->u32Function, pHGCMCall->cParms, rc2));
            *pfLeakIt = true;
        }
        Log(("vbglR0HGCMInternalDoCall: Cancel race ended with rc=%Rrc (rc2=%Rrc) after %llu ms\n",
            rc, rc2, cElapsed));
    }
}

Log(("GstHGCMCall: rc=%Rrc result=%Rrc fu32Flags=%#x fLeakIt=%d\n",
    rc, pHGCMCall->header.result, pHGCMCall->header.fu32Flags, *pfLeakIt));
return rc;
}

```

- VbglGRPerform
- src/VBox/Additions/common/VBoxGuestLib/GenericRequest.cpp

```

DECLVBGL(int) VbglGRPerform (VMMDevRequestHeader *pReq)
{
    RTCCPHYS physaddr;

    // これなんだ
    int rc = vbglR0Enter ();
}

```

```

    if (RT_FAILURE(rc))
        return rc;

    if (!pReq)
        return VERR_INVALID_PARAMETER;

    physaddr = VbglPhysHeapGetPhysAddr (pReq);
    if ( !physaddr
        || (physaddr >> 32) != 0) /* Port IO is 32 bit. */
    {
        rc = VERR_VBGL_INVALID_ADDR;
    }
    else
    {
        ASMOutU32(g_vbgldata.portVMMDev + VMMDEV_PORT_OFF_REQUEST, (uint32_t)physaddr);
        /* Make the compiler aware that the host has changed memory. */
        ASMCompilerBarrier();
        rc = pReq->rc;
    }
    return rc;
}

```

- ASMOutU32
 - include/iprt/asm-amd64-x86.h
 - **ASM = AMD64 and x86 Specific Assembly Functions**
 - outl で I/Oポートにデータを書く
 - VMMDevRequestHeader の開始アドレスを書いて、ホストOSが読む?

```

/**
 * Writes a 32-bit unsigned integer to an I/O port, ordered.
 *
 * @param Port    I/O port to write to.
 * @param u32     32-bit integer to write.
 */
#ifdef RT_INLINE_ASM_EXTERNAL && !RT_INLINE_ASM_USES_INTRIN
DECLASM(void) ASMOutU32(RTIOPORT Port, uint32_t u32);
#else
DECLINLINE(void) ASMOutU32(RTIOPORT Port, uint32_t u32)
{
    # if RT_INLINE_ASM_GNU_STYLE
        __asm__ __volatile__ ("outl %1, %w0\n\t"
                               :: "Nd" (Port),
                               "a" (u32));

    # elif RT_INLINE_ASM_USES_INTRIN
        __outdword(Port, u32);

    # else
        __asm
        {
            mov     dx, [Port]
            mov     eax, [u32]
            out     dx, eax
        }
    # endif
}
#endif

```

- VbglPhysHeapGetPhysAddr
 - src/VBox/Additions/common/VBoxGuestLib/PhysHeap.cpp
 - 物理ヒープ?の物理アドレスを出す?

```

DECLVBGL(uint32_t) VbglPhysHeapGetPhysAddr (void *p)
{
    uint32_t physAddr = 0;
    VBGLPHYSHEAPBLOCK *pBlock = vbglPhysHeapData2Block (p);

    if (pBlock)
    {
        VBGL_PH_ASSERTMsg((pBlock->fu32Flags & VBGL_PH_BF_ALLOCATED) != 0,
                           ("pBlock = %p, pBlock->fu32Flags = %08X\n", pBlock, pBlock->fu32Flags));
    }
}

```



```

        if (pBlock->fu32Flags & VBGL_PH_BF_ALLOCATED)
            physAddr = pBlock->pChunk->physAddr + (uint32_t)((uintptr_t)p - (uintptr_t)pBlock->pChunk);
    }

    return physAddr;
}

```

HostServices

ホストOS側のVirtualBoxVM プロセスで実行されるユーザランドのコード。Not Kernel

- src/VBox/HostServices/SharedFolders/vbsf.h に ホストOSが svcCall で呼び出すAPIが定義されている
- ホストOSの VirtualBoxVMプロセスが呼び出すシステムコールの抽象化ラッパー
 - ホストOSのプラットフォームによって呼び出すべきシステムコールが違うので抽象化をかます必要がある

```

// ...

int vbsfCreate (SHFLCLIENTDATA *pClient, SHFLROOT root, SHFLSTRING *pPath, uint32_t cbPath, SHFLCREATEPARAMS *pParams);

int vbsfClose (SHFLCLIENTDATA *pClient, SHFLROOT root, SHFLHANDLE Handle);

int vbsfRead (SHFLCLIENTDATA *pClient, SHFLROOT root, SHFLHANDLE Handle, uint64_t offset, uint32_t *pcbBuffer, uint32_t *pError);
int vbsfWrite (SHFLCLIENTDATA *pClient, SHFLROOT root, SHFLHANDLE Handle, uint64_t offset, uint32_t *pcbBuffer, uint32_t *pError);
int vbsfLock (SHFLCLIENTDATA *pClient, SHFLROOT root, SHFLHANDLE Handle, uint64_t offset, uint64_t length, uint32_t *pError);
int vbsfUnlock (SHFLCLIENTDATA *pClient, SHFLROOT root, SHFLHANDLE Handle, uint64_t offset, uint64_t length, uint32_t *pError);
int vbsfRemove (SHFLCLIENTDATA *pClient, SHFLROOT root, SHFLSTRING *pPath, uint32_t cbPath, uint32_t flags);
int vbsfRename (SHFLCLIENTDATA *pClient, SHFLROOT root, SHFLSTRING *pSrc, SHFLSTRING *pDest, uint32_t flags);
int vbsfDirList (SHFLCLIENTDATA *pClient, SHFLROOT root, SHFLHANDLE Handle, SHFLSTRING *pPath, uint32_t flags, uint32_t *pError);
int vbsfFileInfo (SHFLCLIENTDATA *pClient, SHFLROOT root, SHFLHANDLE Handle, uint32_t flags, uint32_t *pcbFileInfo, uint32_t *pError);
int vbsfQueryFSInfo (SHFLCLIENTDATA *pClient, SHFLROOT root, SHFLHANDLE Handle, uint32_t flags, uint32_t *pQueryFSInfo, uint32_t *pError);
int vbsfSetFSInfo (SHFLCLIENTDATA *pClient, SHFLROOT root, SHFLHANDLE Handle, uint32_t flags, uint32_t *pSetFSInfo, uint32_t *pError);
int vbsfFlush (SHFLCLIENTDATA *pClient, SHFLROOT root, SHFLHANDLE Handle);
int vbsfDisconnect (SHFLCLIENTDATA *pClient);
int vbsfQueryFileInfo (SHFLCLIENTDATA *pClient, SHFLROOT root, SHFLHANDLE Handle, uint32_t flags, uint32_t *pQueryFileInfo, uint32_t *pError);
int vbsfReadLink (SHFLCLIENTDATA *pClient, SHFLROOT root, SHFLSTRING *pPath, uint32_t cbPath, uint8_t *pBuf, uint32_t *pError);
int vbsfSymlink (SHFLCLIENTDATA *pClient, SHFLROOT root, SHFLSTRING *pNewPath, SHFLSTRING *pOldPath, SHFLSTRING *pError);

#endif /* __VBSF__H */

```

vbsfCreate

ホストOSがUNIX系OSなら、最終的に open(2) を呼び出すはず

- vbfsOpenFile か vbsfOpenDir に続く
 - vbsfOpenFile から RTFileOpen を呼び出す
 - **RT** = RunTime ?

```

// include/iprt/mangling.h
# define RTFileOpen                                RT_MANGLER(RTFileOpen)

// include/VBox/VBoxGuestMangling.h
#define RT_MANGLER(symbol)    VBoxGuest_##symbol

// include/VBox/SUPDrvMangling.h
#define RT_MANGLER(symbol)    VBoxHost_##symbol

// 最終的に VBoxGuest_RTFileOpen, VBoxH0ST_RTFileOpen のシンボルに変換される?

```

- src/VBox/Runtime 以下にプラットフォームごとのディレクトリが見つかる
 - src/VBox/Runtime/r3/posix/fileio-posix.cpp に POSIX な RTFileOpen実装がある
 - (ホストOSのプロセスとして) open(2) を呼び出している

```

RTR3DECL(int) RTFileOpen(PRTFILE pFile, const char *pszFilename, uint64_t fOpen)
{
    // ....

    /*
     * Open/create the file.
     */
}

```

```
*/  
char const *pszNativeFilename;  
rc = rtPathToNative(&pszNativeFilename, pszFilename, NULL);  
if (RT_FAILURE(rc))  
    return (rc);  
  
int fh = open(pszNativeFilename, fOpenMode, fMode);  
int iErr = errno;
```

ということで RT_MANGLER を見つけたら src/VBox/Runtime/ 以下の実装を追えばok

IPRT って何?

IPRT, a portable runtime library which abstracts file access, threading, string manipulation, etc.
Whenever VirtualBox accesses host operating features, it does so through this library for cross-platform po

<http://www.virtualbox.org/manual/ch10.html>

